

Transactional systems

COURSE 9: Databases

Transactional systems

Transaction

- Set of operations on the database, set of statements:
 - insert, update, delete
- Delimited by statements or function calls:
 - begin transaction
 - end transaction
- All operations are finalized with success, or none is saved in the db.
- A transactional system must
 - manage concurrent transactions.
 - ensure consistent data in case of failure.

Transaction

Statement 1

Statement 2

commit -- end transaction 1

Statement 3

Statement 4

Statement 5

commit -- end transaction 2

Transaction properties

ACID



ATOMICITY

CONSISTENCY

ISOLATION

DURABILITY

- all changes must be saved
 - collection of steps → single indivisible unit.
- If one operation fails all changes to the database must be undone
 - Failures in transaction, example: statement error, violating unique constraints.
 - System failures, OS crash.




ATOMICITY

CONSISTENCY

ISOLATION

DURABILITY

- If a transaction starts in a consistent state, the database must be consistent at the end of the transaction.
 - Database constraints
 - PRIMARY KEY key constraint,
 - UNIQUE,
 - NOT NULL,
 - FOREIGN KEY referential integrity,
 - CHECK
 - Business constraints - triggers.



- The database may at some point be in an *inconsistent state*.
- Inconsistencies are not visible in a database system (ensured by *atomicity*).
- The old values of any data on which a transaction performs is written to a log file used by a
→ *recovery system*.

ATOMICITY

CONSISTENCY

ISOLATION

DURABILITY

- The database system must ensure that transactions run without interference.
 - For any pair of transactions T_i, T_j ,
first statement of transaction T_i is executed after T_j finished or
first statement of transaction T_j is executed after T_i finished.



ATOMICITY

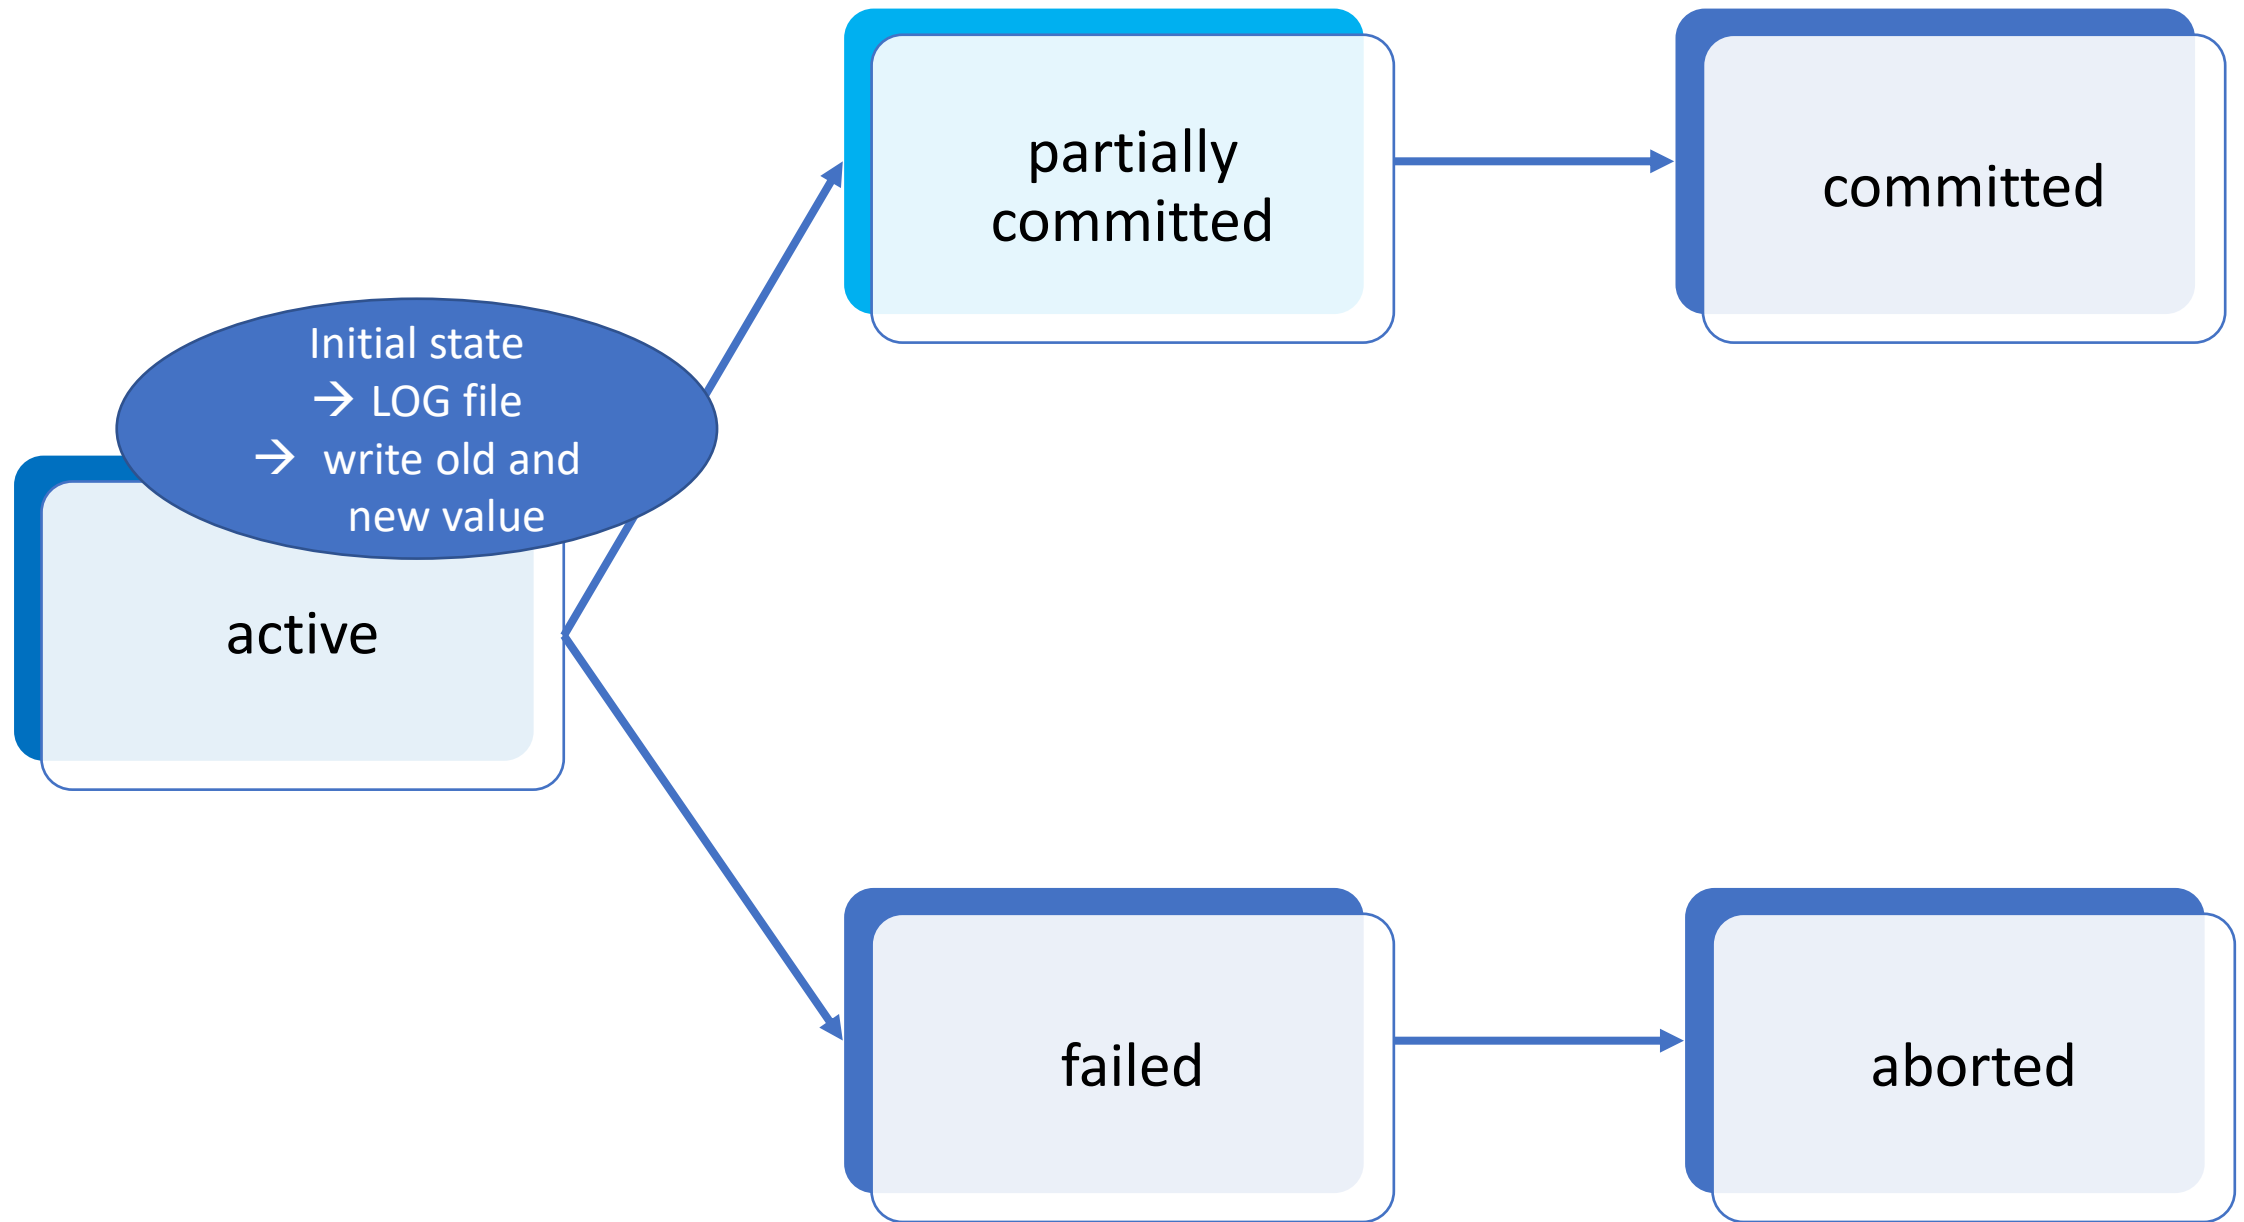
CONSISTENCY

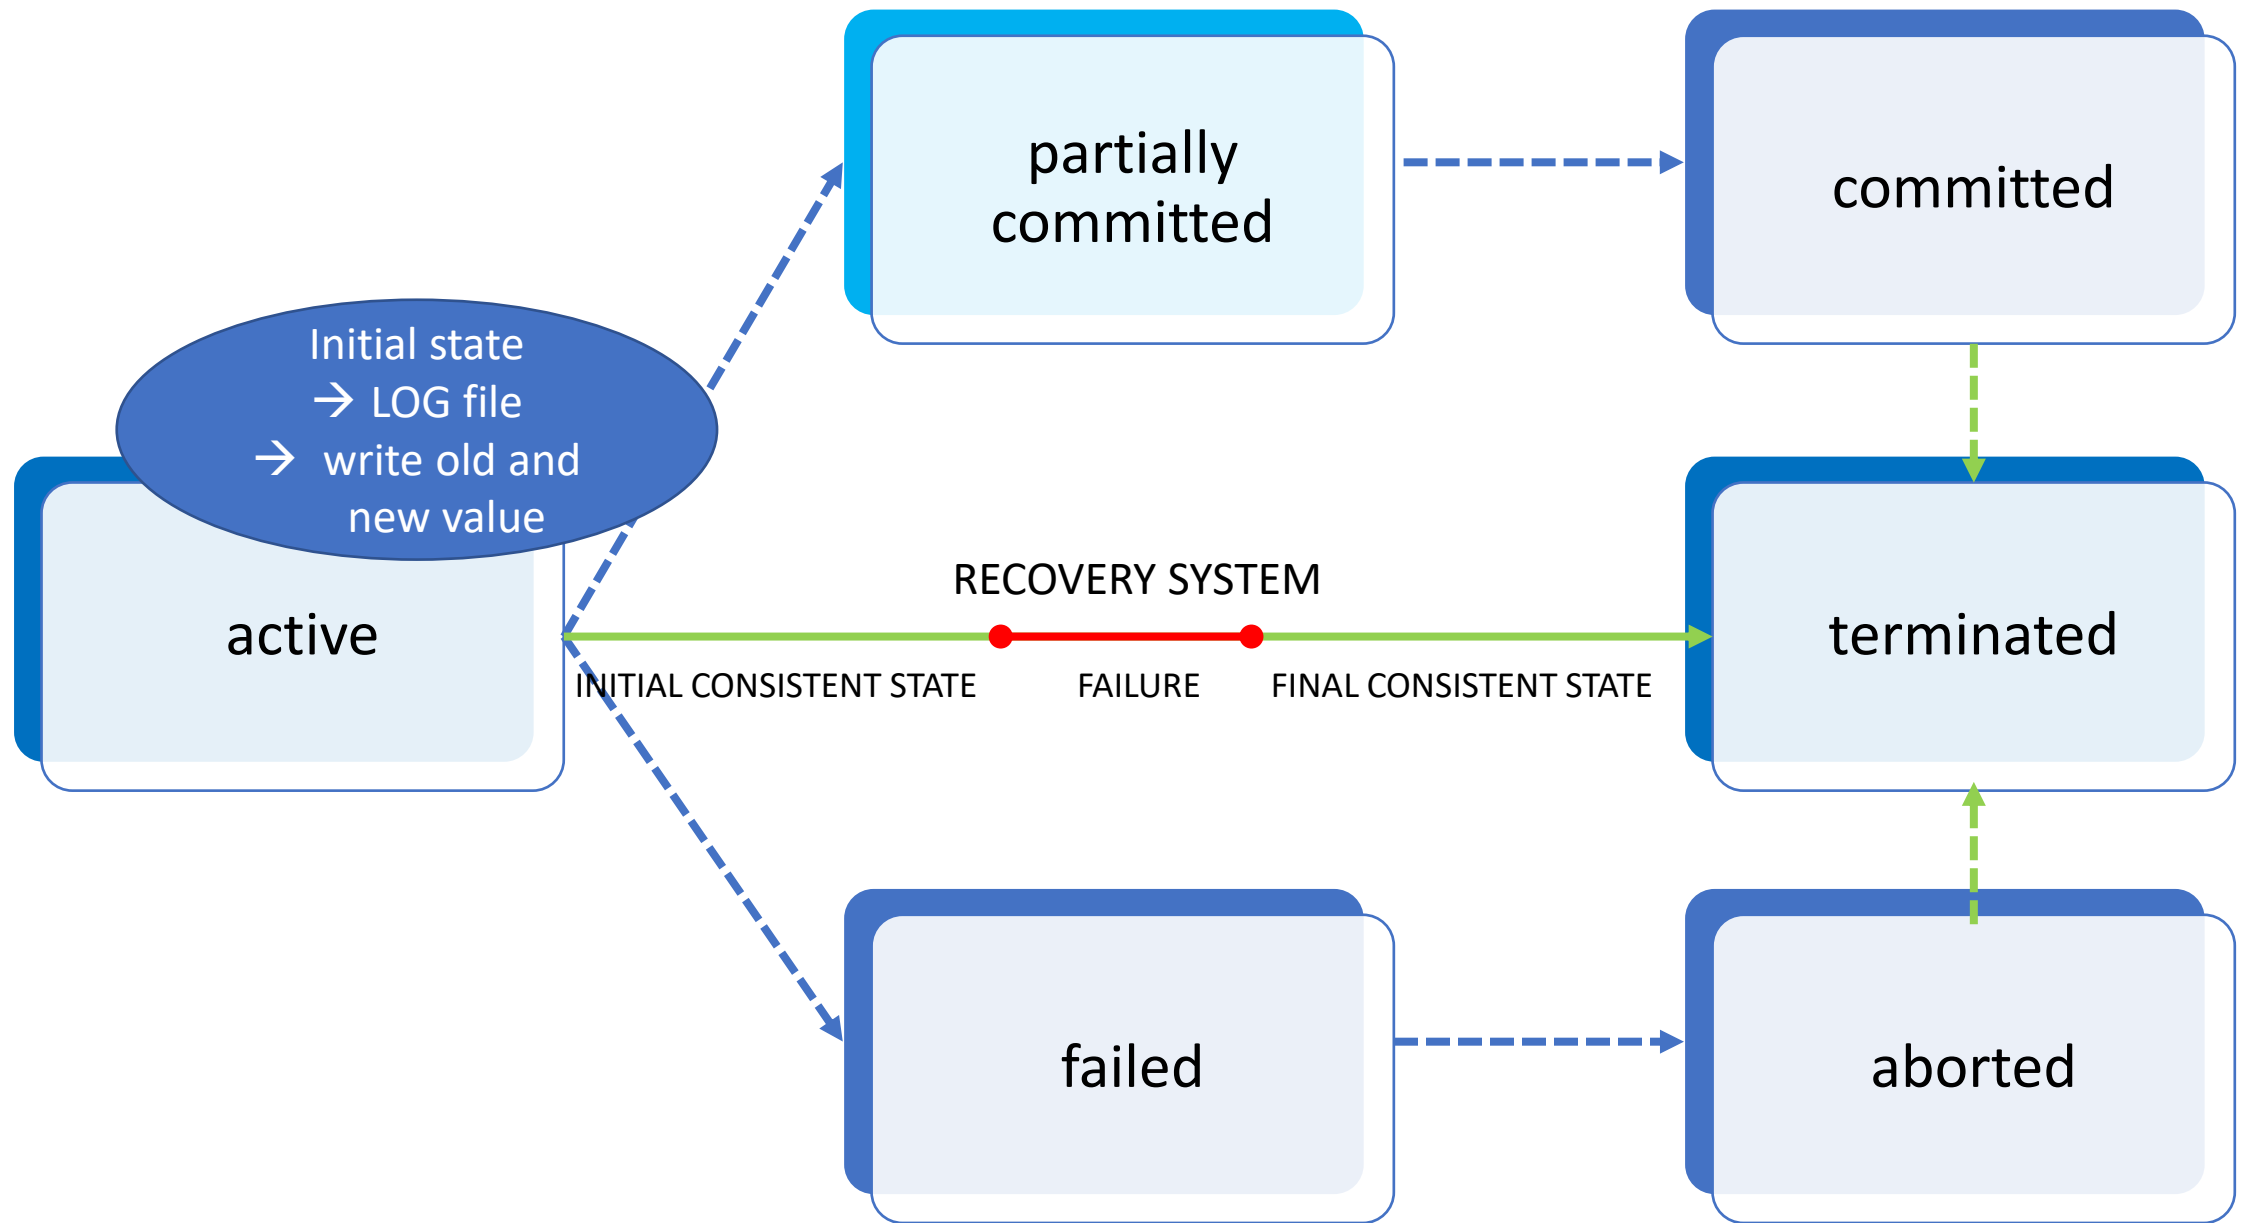
ISOLATION

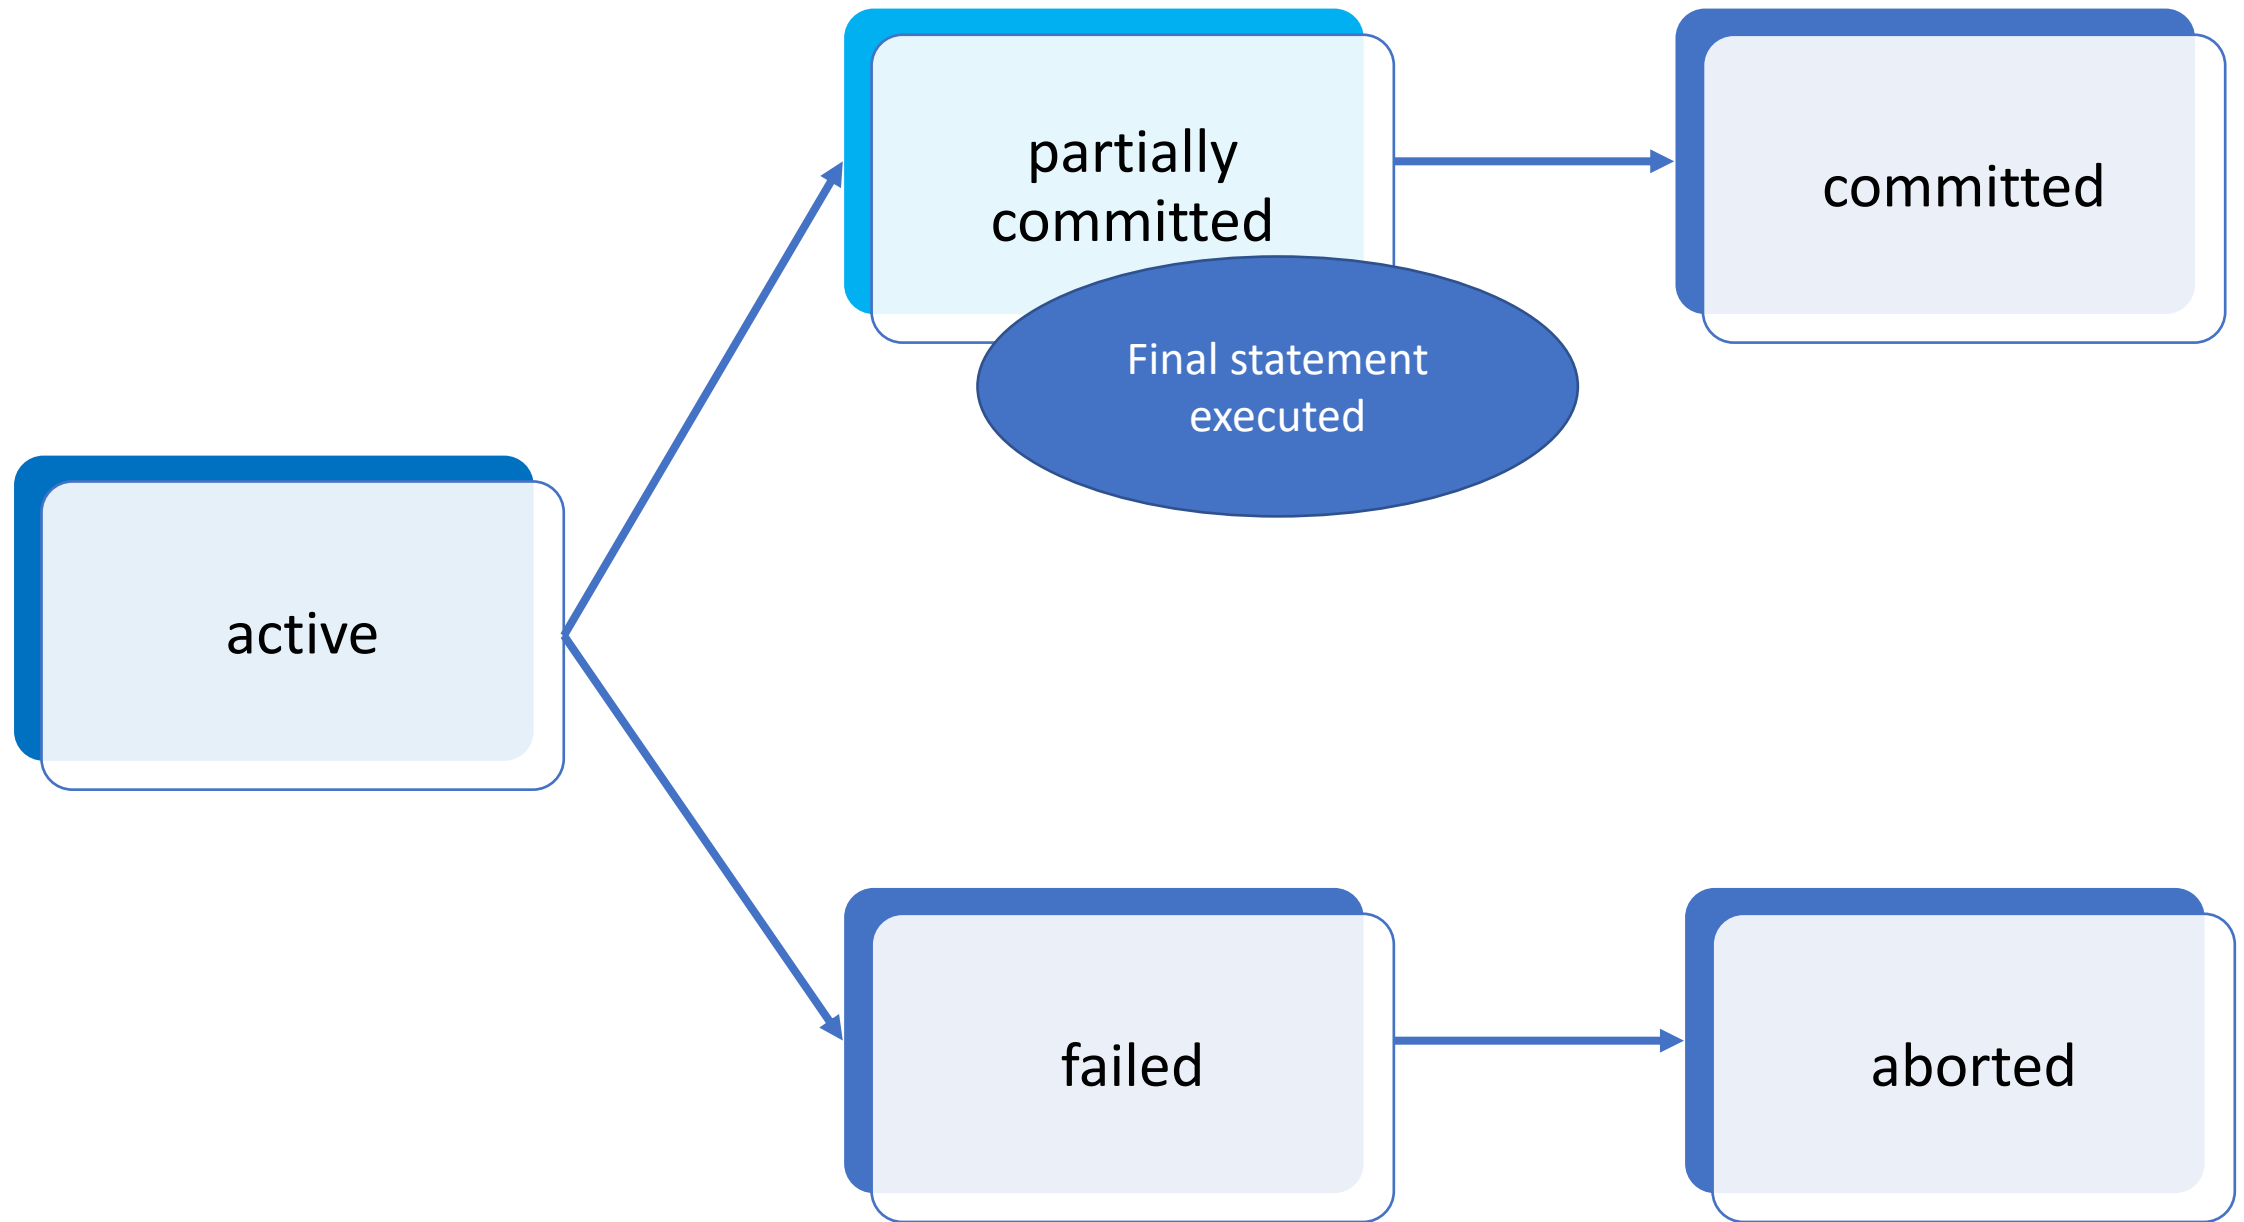
DURABILITY

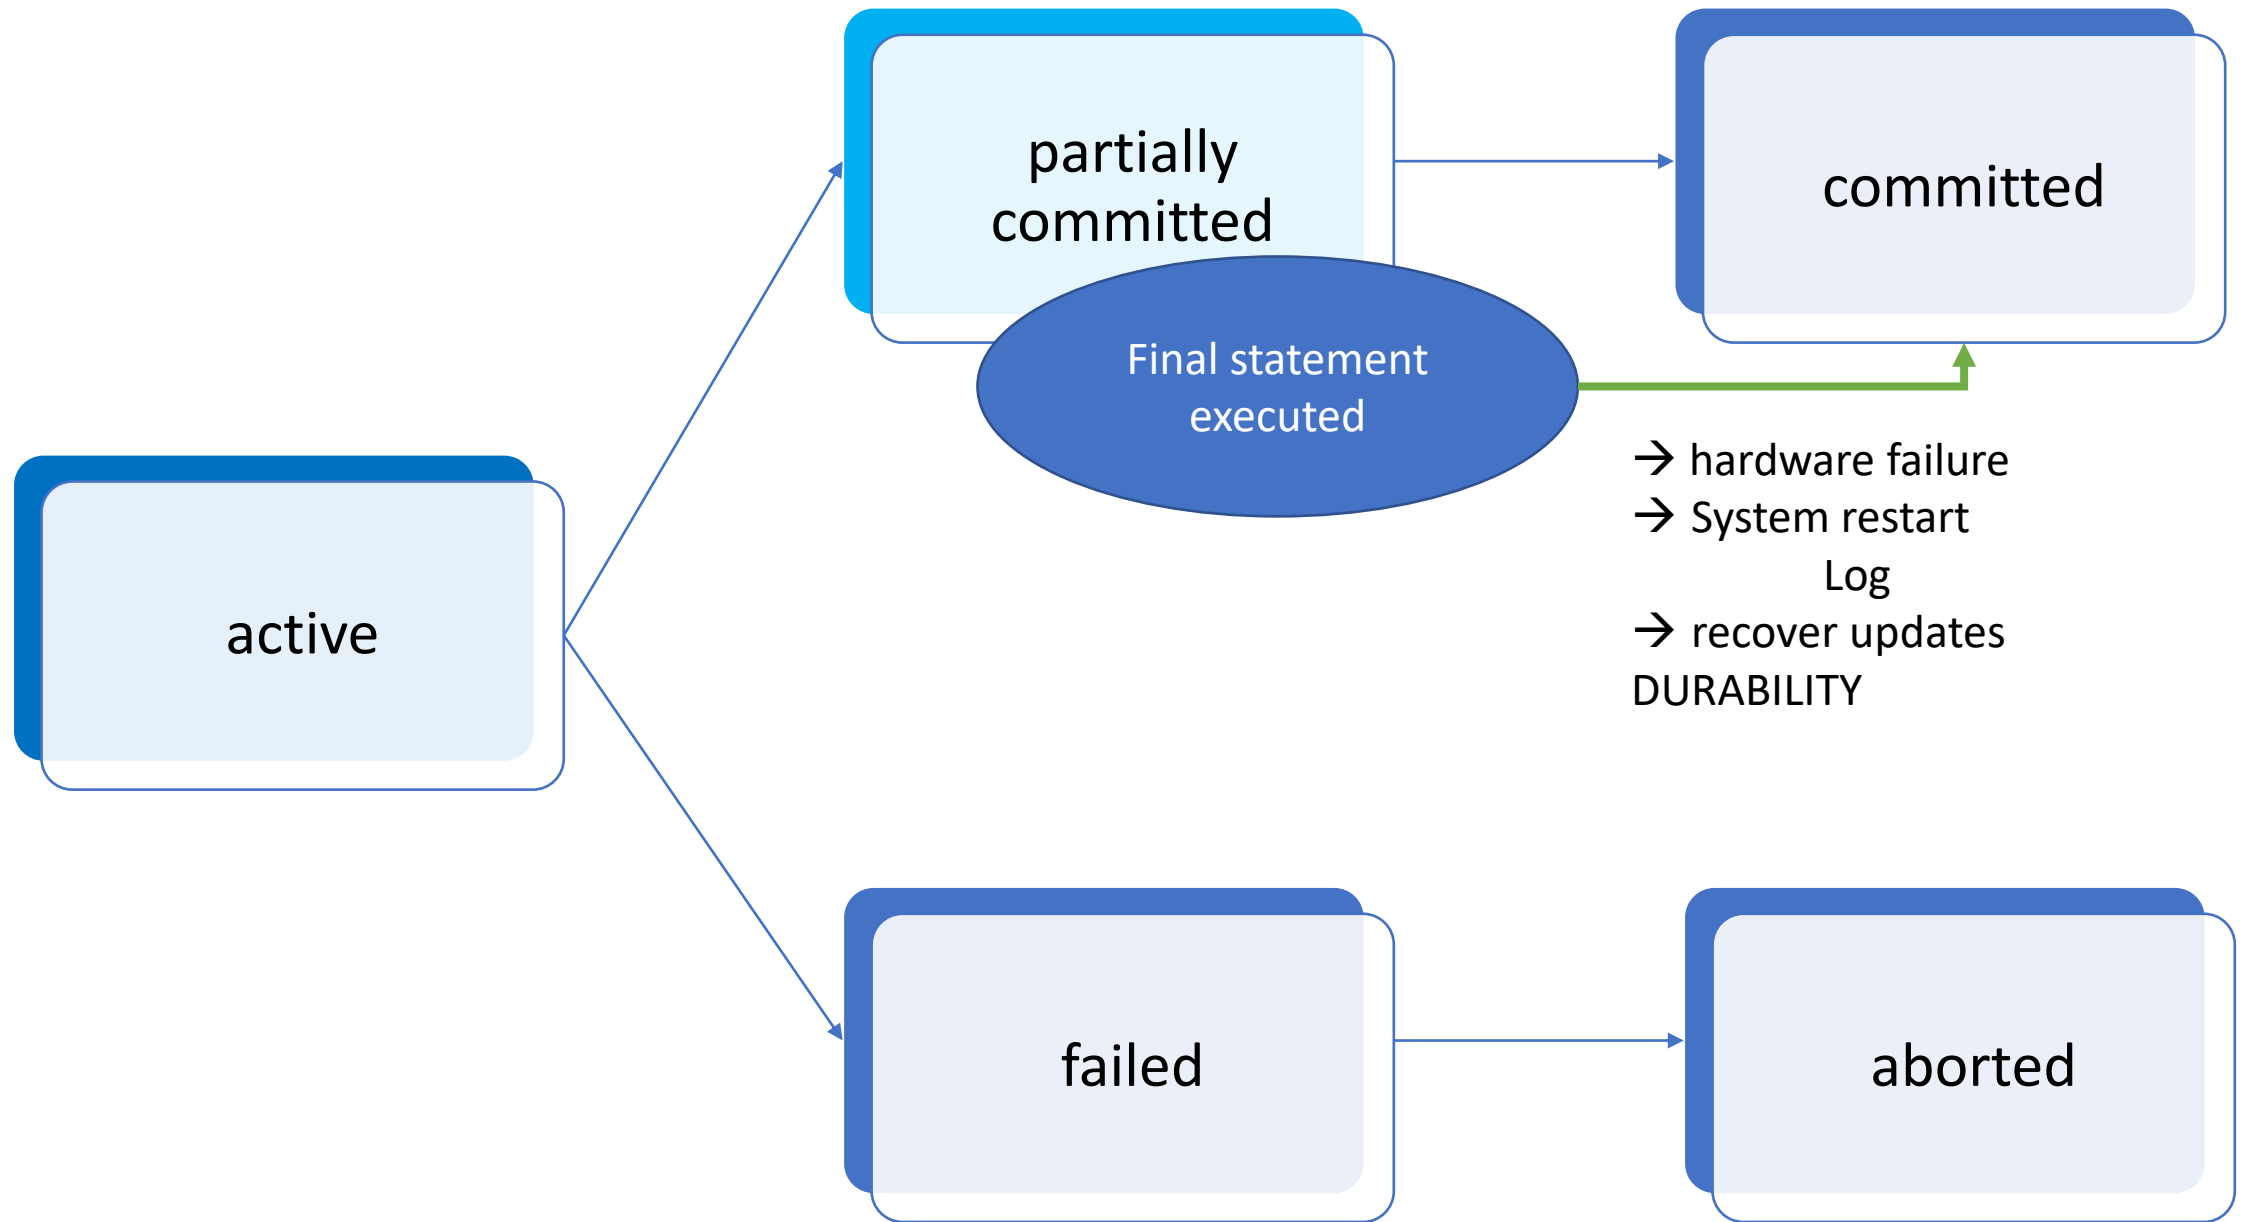
- After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.
- Information about the updates performed by the transaction is written to disk and used to reconstruct the database after failure.
→ *recovery system.*


Transaction states

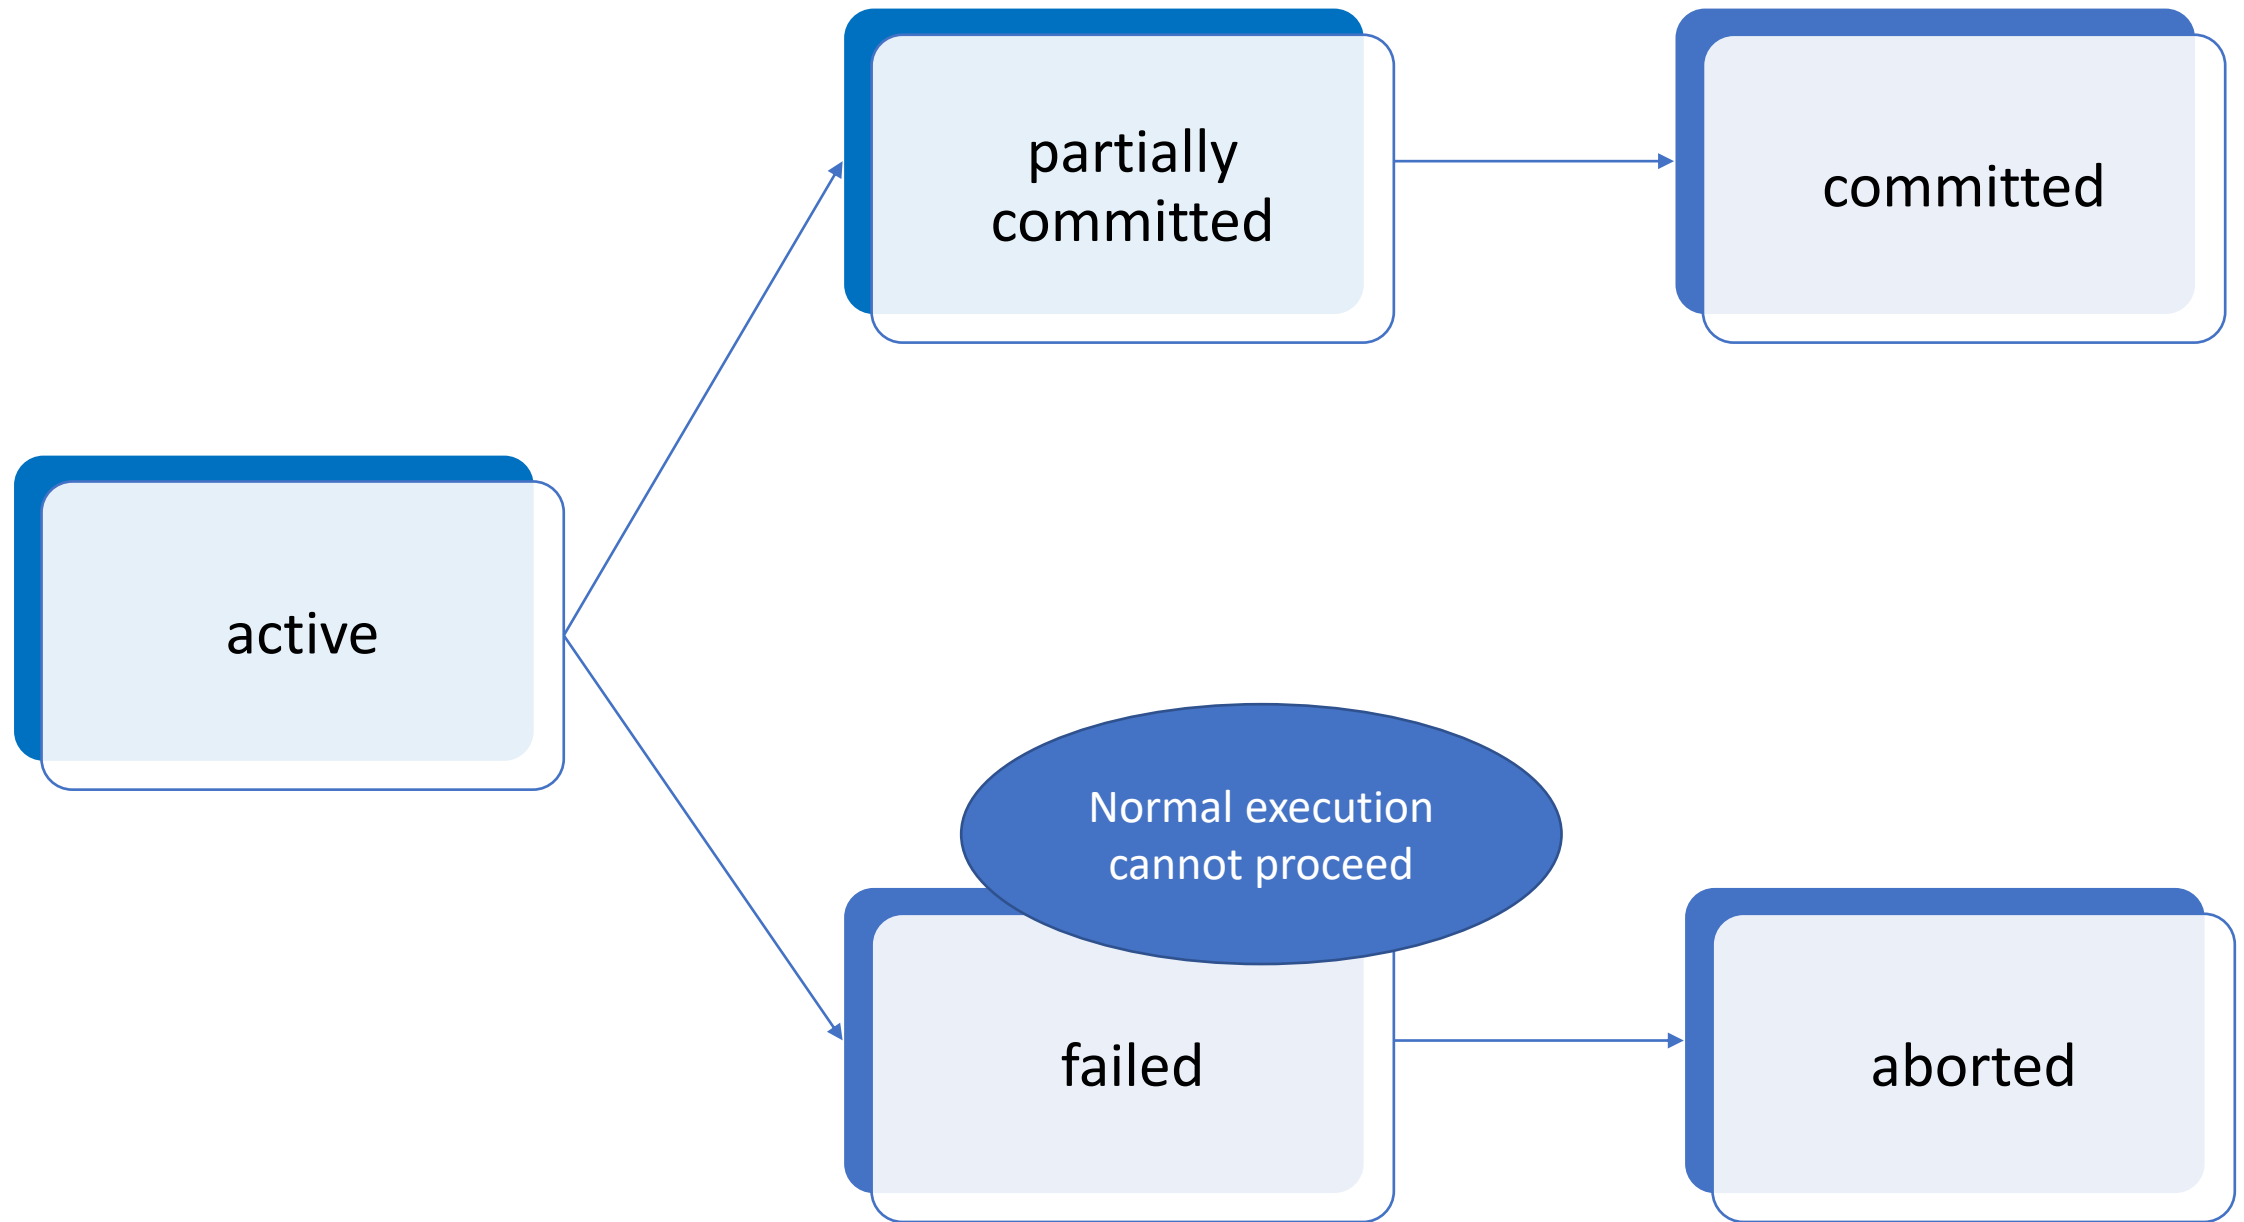


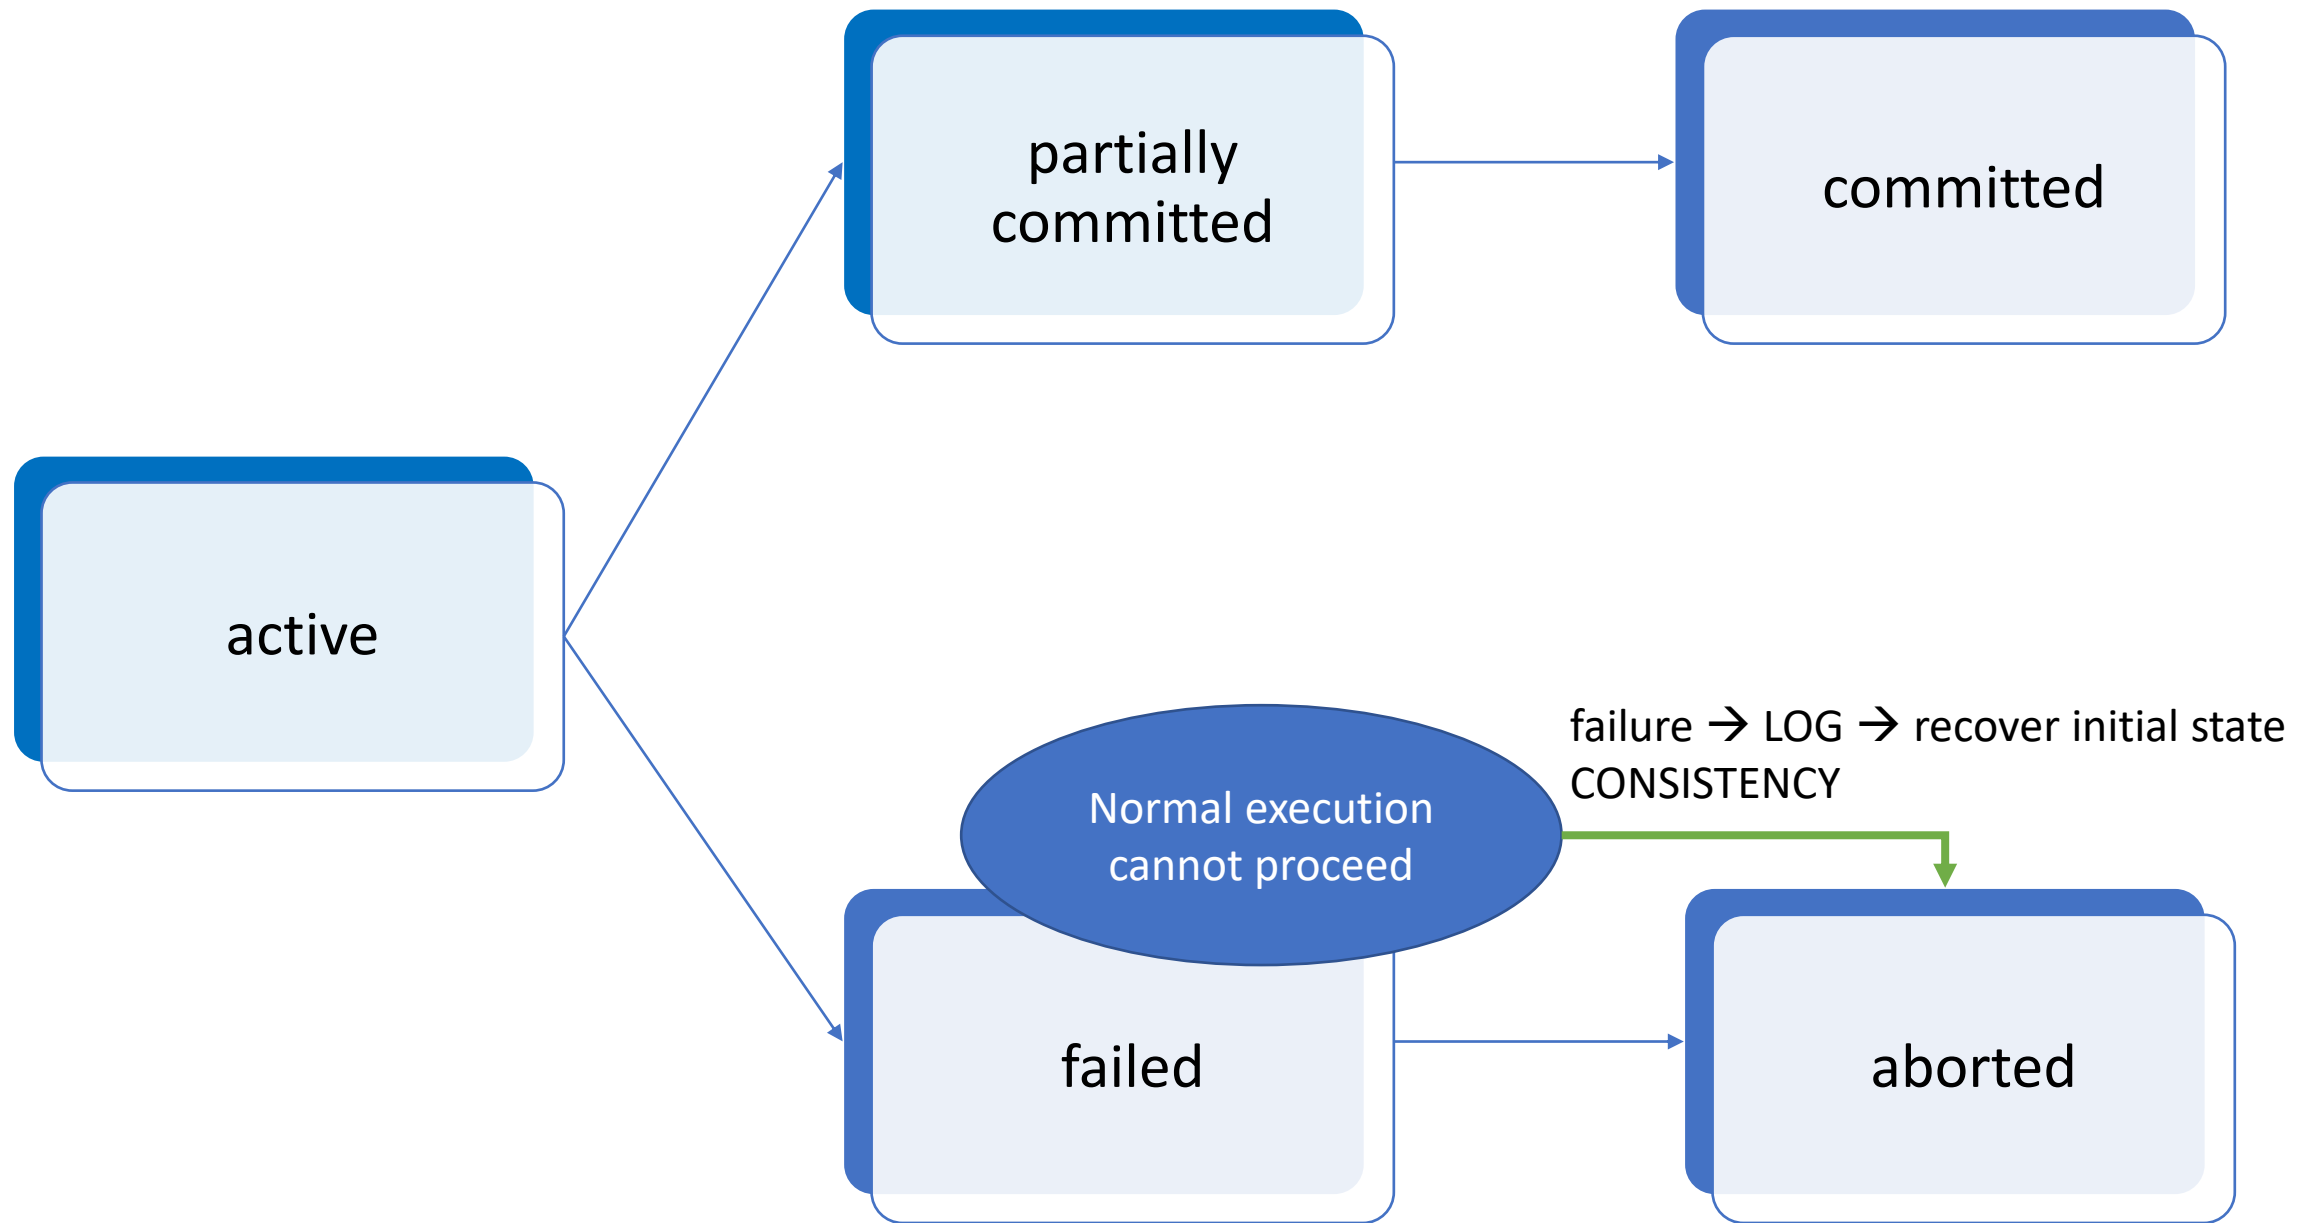


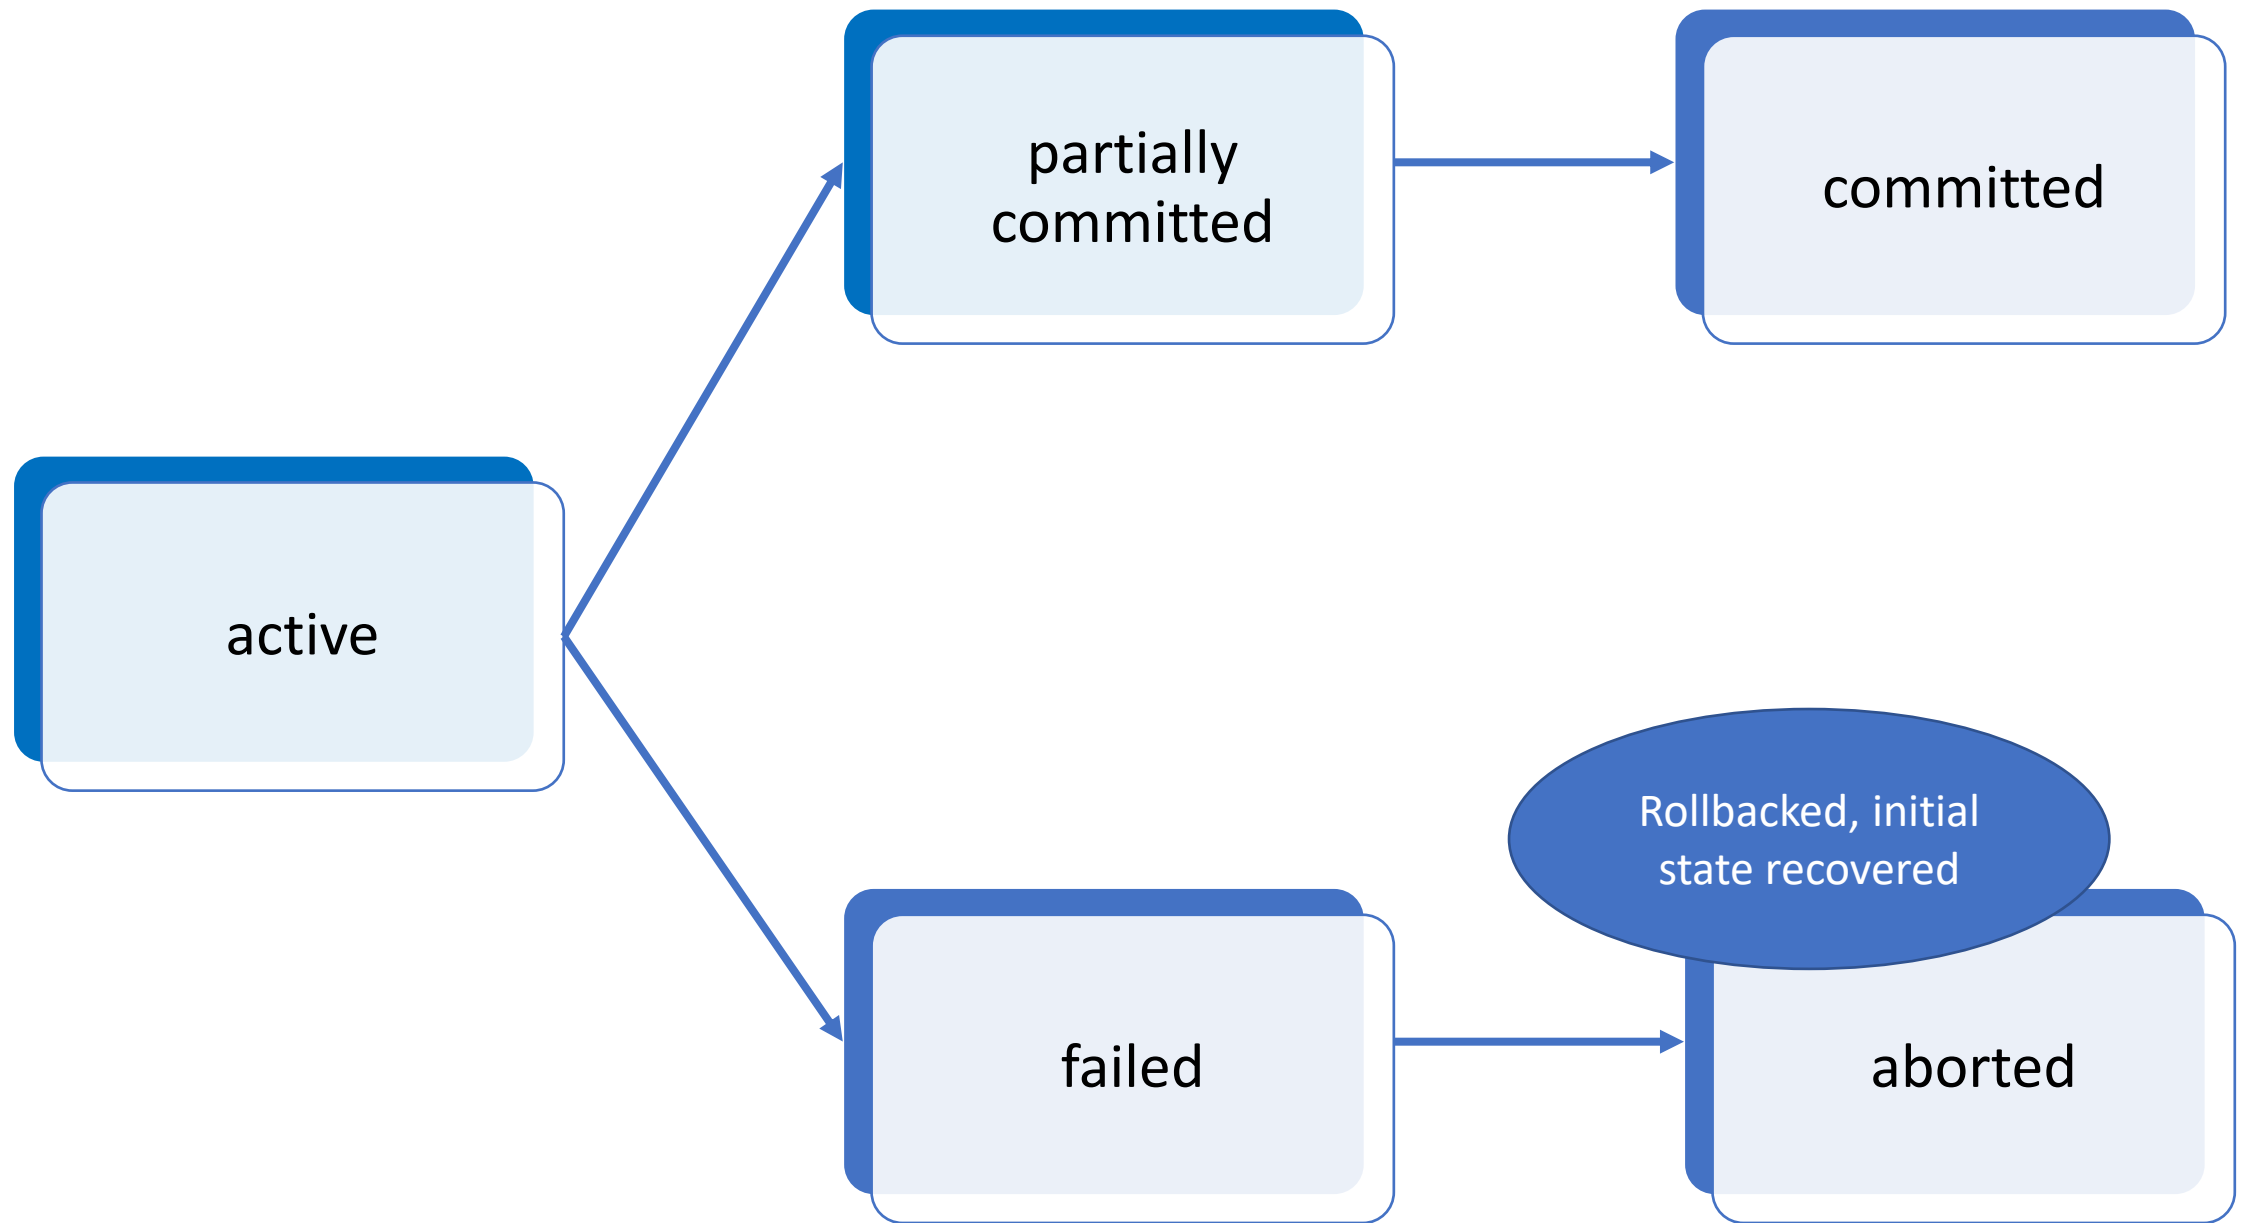












Concurrent transactions

Concurrent transactions

- Reduce response time (time for a transaction to be completed).
- Improved workload/resource utilization.
- ISOLATION may be violated, as a result database may be found in an inconsistent state
 - *Concurrency control*

Concurrent transactions - *conflicts*

- Serial execution preserves consistency, assuming that each transactions preserve consistency.

first statement of transaction T_i is executed after T_j finished or
first statement of transaction T_j is executed after T_i finished

single threaded transactions.

- Instructions I of T_i and J of T_j conflict \Leftrightarrow there exists a *data* accessed by both I and J , and at least one of I and J write *data*.

1. $I = \text{read}(\text{data})$	$J = \text{read}(\text{data})$	I and J not conflicting.
2. $I = \text{read}(\text{data})$	$J = \text{write}(\text{data})$	conflict
3. $I = \text{write}(\text{data})$	$J = \text{read}(\text{data})$	conflict
4. $I = \text{write}(\text{data})$	$J = \text{write}(\text{data})$	conflict

Concurrent transactions -- Schedules

- **Schedules:** sequences of instructions that specify the chronological order in which instructions of concurrent transactions are executed
 - A schedule for a set of transactions must consist of **all instructions of all transactions**.
 - A schedule must **preserve the order** in which the instructions appear in each individual transaction.
 - last statement commit or abort.

Schedules example S1

- Serial execution.
- No conflicts.
- DB in consistent state

$$A.\text{new} + B.\text{new} = A.\text{old} + B.\text{old}$$

T1	T2
<pre>read (A) A := A - 50 write (A) read (B) B := B + 50 write (B) commit</pre>	
	<pre>read (A) temp := A * 0.1 A := A - temp write (A) read (B) B := B + temp write (B) commit</pre>

Schedules example S2

- Not a serial execution.
- Equivalent to Schedule S1.
- DB in consistent state

$$A.\text{new} + B.\text{new} = A.\text{old} + B.\text{old}$$

T1	T2
read (A) A := A - 50 write (A)	
	read (A) temp := A * 0.1 A := A - temp write (A)
read (B) B := B + 50 write (B) commit	
	read (B) B := B + temp write (B) commit

Concurrent transactions

- A (possibly concurrent) schedule is serializable if it is equivalent to a serial schedule. Different forms of schedule equivalence:
 1. Conflict serializability
 2. View serializability

Concurrent transactions

- A (possibly concurrent) schedule is serializable if it is equivalent to a serial schedule. Different forms of schedule equivalence:
 1. Conflict serializability

If a schedule S can be transformed into a schedule S' by a series of swaps of non-conflicting instructions, we say that S and S' are conflict equivalent.
 2. View serializability

Schedules example S2

- Not a serial execution.
- Equivalent to Schedule S1.
- DB in consistent state
 - $A_{\text{new}} + B_{\text{new}} = A_{\text{old}} + B_{\text{old}}$

not conflicting.
by swapping the two
blocks we obtain S1

T1	T2
read (A) A := A - 50 write (A)	
	read (A) temp := A * 0.1 A := A - temp write (A)
read (B) B := B + 50 write (B) commit	
	read (B) B := B + temp write (B) commit

Schedules example S3

- Not a serial execution.
- Not equivalent to Schedule S1.
- DB in inconsistent state
 - $A_{\text{new}} + B_{\text{new}} \neq A_{\text{old}} + B_{\text{old}}$

conflicting,
A is updated by both
blocks

T1	T2
read (A) A := A - 50	
	read (A) temp := A * 0.1 A := A - temp write (A)
write (A) read (B) B := B + 50 write (B) commit	
	read (B) B := B + temp write (B) commit

Concurrent transactions

1. Conflict serializability
2. View serializability

Let S and S' be 2 schedules with the same set of transactions. S and S' are view equivalent if the following 3 conditions are met, for each data item Q :

- **initial value:** If in schedule S , transaction T_i reads the **initial value of Q** , then in schedule S' also transaction T_i must read the initial value of Q .
- **write-read order:** If in schedule S transaction T_i executes $\text{read}(Q)$, and that value was produced by transaction T_j (if any), then in schedule S' transaction T_i must read the value of Q that was produced by the same $\text{write}(Q)$ operation of transaction T_j .
- **final value:** The transaction (if any) that performs the **final $\text{write}(Q)$** operation in schedule S must also perform the final $\text{write}(Q)$ operation in schedule S' .

View equivalence is based purely on reads and writes alone.

Concurrent transactions

- Test serializability :


2. View serializability

- The problem of checking if a schedule is view-serializable falls in the class of NP-complete problems. Thus, existence of an efficient algorithm is extremely unlikely.
- Practical algorithms that just check some sufficient conditions for view serializability can still be used.

Concurrent transactions

- Test serializability:
 1. Conflict serializability
 - Consider some schedule of a set of transactions T_1, T_2, \dots, T_n
 - Precedence graph — a direct graph where the vertices are the transactions (names).
 - We draw an arc from T_i to T_j if the second transaction conflicts, and T_i accessed the conflicting data item.
 - We may label the arc by the data item that was accessed.
 - A schedule is CS if and only if its precedence graph is acyclic.
 - If precedence graph is acyclic, the serializability order can be obtained by a **topological sorting** of the graph.

Conflict serializability



conflicting,
A is updated by both
blocks

T1: read (A) A := A - 50	
	T2: read (A) temp := A * 0.1 A := A - temp write (A)
T1: write (A) read (B) B := B + 50 write (B) commit	
	T3: read (B) B := B + temp write (B) commit

Isolation levels

Isolation levels

- **Isolation:** execute a transaction *as if* there are no other concurrent transactions running simultaneously.
 - Prevent read or write of incorrect, temporary, aborted data processed by concurrent transactions
- **Isolation levels:** trade off between *perfect* isolation and performance
 - response time: time elapsed before a transaction completes
 - throughput: number of transactions per second

Level **Serializability**, perfect isolation

- The final state of the database is equivalent to a state of the database if the transactions were run sequentially.
 - serializable schedule.
- Way of obtaining serializability:
 - locking.
 - timestamp validation.
 - multi-versioning.

Read phenomenon

lost-update anomaly

final stock 12!

T1	T2
<pre>select qte into :nS from stock where n_prod = 100 --nS = 13</pre>	
	<pre>select qte into :nS from stock where n_prod = 100</pre>
<pre>update stock set qte = :nS - 2 where n_prod = 100</pre>	
	<pre>update stock set qte = :nS - 1 where n_prod = 100</pre>
<pre>insert into orders(n_prod, qte) values(100, 2) commit</pre>	
	<pre>insert into orders(n_prod, qte) values(100, 1)</pre>

Read phenomenon

dirty-read anomaly
number of products
(in/out) = qte_stock
1 product missing!
Read uncommitted
data

T1	T2
<pre>select qte into :nS from stock where n_prod = 100 --nS = 13</pre>	
<pre>update stock set qte = :nS - 1 where n_prod = 100</pre>	
	<pre>select sum(qte) into :nM from mvt where n_prod = 100</pre>
	<pre>select qte into :nS from stock where n_prod = 100</pre>
<pre>insert into mvt(n_prod, qte) values(100, -1) commit</pre>	

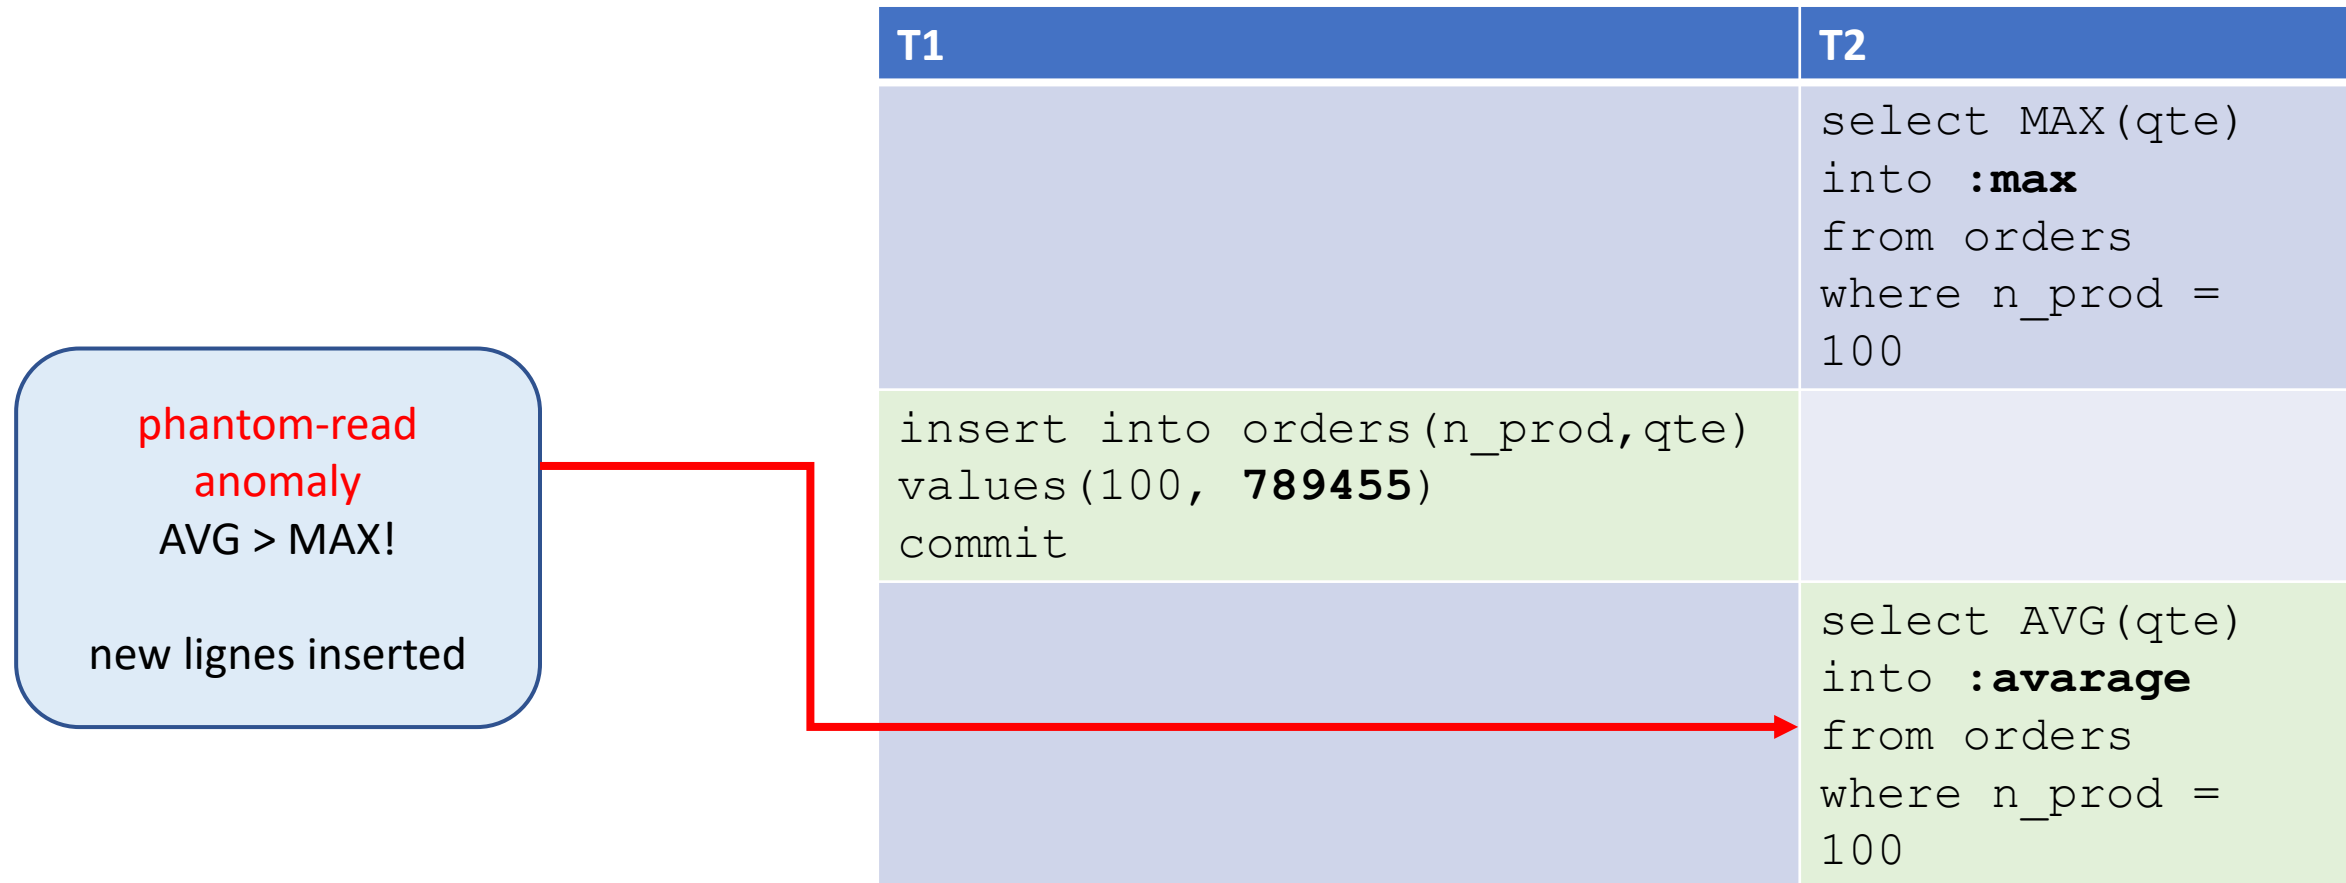
Read phenomenon

non-repeatable read
anomaly

only one insert into
restock is needed!
read twice, different
values

T1	T2
...	<pre>select qte into :nS from stock where n_prod = 100 --nS = 10</pre>
	<pre>if nS < 15 and nS >= 10 insert into restock(n_prod, qte) values(100, 5)</pre>
<pre>update stock set qte = :nS - 1 where n_prod = 100 insert into orders (n_prod, qte) values(100, 1) commit</pre>	
	<pre>select qte into :nS from stock where n_prod = 100 if nS < 10 insert into restock(n_prod, qte) values(100, 15)</pre>

Read phenomenon







Isolation levels

- weaker the isolation level → more anomalies may occur





	ERROR	dirty-reads	non-repeatable reads	phantom
LEVEL				
READ UNCOMMITTED		✓	✓	✓
READ COMMITTED		✗	✓	✓
REPEATABLE READ		✗	✗	✓
SERIALIZABLE		✗	✗	✗

Isolation levels

	ERROR	lost-update	dirty-reads	non-repeatable reads	phantom
LEVEL					
READ UNCOMMITTED					





- allows uncommitted data to be read
- all isolation levels prevent writes to a data item that has already been written by another transaction not yet committed or aborted (rollbacked).

Isolation levels

	ERROR	lost-update	dirty-reads	non-repeatable reads	phantom
LEVEL					
READ COMMITTED					




- read only committed
- does not require repeatable reads. Between two reads of a data item by the transaction, another transaction may have updated the data item and committed.

Isolation levels

	ERROR	lost-update	dirty-reads	non-repeatable reads	phantom
LEVEL					
REPEATABLE READ					

- read only committed
- between two reads of an item by a transaction, no other transaction is allowed to update it.
- a transaction may find other data inserted by a committed transaction

Isolation levels

	ERROR	lost-update	dirty-reads	non-repeatable reads	phantom
LEVEL					
SERIALIZABLE					

- read only committed
- between two reads of an item by a transaction, no other transaction is allowed to update it.
- a transaction may find other data inserted by a committed transaction

Achieving isolation

- Versioning
 - Transactions read from a “snapshot” of the database (timestamp-versioning)
- Locking
 - Read or write locks

Locking

- Locks prevent destructive interactions between transactions accessing the same resource.
 - Shared access to read
 - Exclusive access to read and write
- Locks (Shared, Shared) compatible.
- Locks (Shared, Exclusive) not compatible.
- A transaction waits until all incompatible locks held by other transactions are released.
- <https://oracle-base.com/articles/misc/deadlocks>
- https://docs.oracle.com/cd/B19306_01/server.102/b14220/consist.htm

Snapshot isolation

- Snapshot of the database at the beginning of each transaction.
- The transaction operates only on that snapshot.
- The snapshot consists only of committed values.
- Updates are kept in transaction workspace until commit.
- Implemented with timestamp-versioning

BASE

NoSql consistency model